

Sanity Checks for Reinforcement Learning Explanations

Aaryan Singh

Independent Researcher

singha9@rose-hulman.edu

Abstract. Interpretability is proposed as a basis for oversight: if we can see which inputs drive a policy, we can decide whether to trust it. That argument needs attributions to carry information about the agent, and we show they need not. On a prediction market that is empirically efficient, a PPO agent whose return is statistically indistinguishable from abstention produces attributions with a cross-seed rank correlation of 0.849, unanimous agreement on the most important feature, and a top-5 ranking whose adjacent gaps exceed their own standard errors. Every conventional check passes and the explanation is empty. We introduce a null-model test that detects this, built on the additivity identity attribution schemes already satisfy: compare an agent’s attribution span against the spans of agents trained where the observation channel carries no information. It rejects a planted signal at $z = 12.27$, fails to reject the market agent at $z = +0.23$, and fires at $z = +3.40$ on a genuinely learnable task built from the same real episodes, so it separates tasks by whether their observations carry usable information rather than by whether the corpus is synthetic. Two results then concern practice. Explanations are cheap to forge: an auxiliary training penalty drives a feature’s attribution from 39.9% to 3.2% of total mass at no measurable cost in return, with no deception and no attack on the explainer, so an oversight regime reading attributions as evidence of what a model relies on is reading a quantity the developer can set wherever the feature does not matter. And the established sanity check for this purpose, randomising network parameters, does not transfer to reinforcement learning: across sixteen checkpoints on four control tasks it never exceeds $z = 2.45$, because its reference distribution equals the spread of random-initialisation return to three significant figures and therefore measures the variance of initialisation, not the explanation.

1 Introduction

A common argument for interpretability is oversight: if we can see which inputs drive a policy, we can decide whether to trust it. This requires that an attribution mean something. It is not obvious that attributions do, and it is less obvious how one would tell.

The difficulty is that attribution methods return an answer regardless. Given any policy and any state, Shapley values [8, 11], integrated gradients [13], and permutation importance all produce a vector of numbers, ranked, signed, and plottable. Nothing in the output distinguishes an attribution reflecting structure the agent exploits from one reflecting nothing at all.

In supervised learning this is handled by randomisation tests. Adebayo et al. [1] compare explanations of a trained model against explanations of a model with randomised parameters, and against a model trained on permuted labels. If the explanation does not change, it does not depend on what the model learned. These tests are standard, and applying them invalidated several saliency methods.

Reinforcement learning has adopted only half of this. Huber et al. [7] apply the parameter-randomisation test to saliency maps for Atari agents, and observe that the data-randomisation test has no obvious RL analogue: there is no labelled dataset to permute, and it is unclear how a sequential decision maker should respond to artificially corrupted features.

This paper supplies that analogue, and shows that the

half already adopted does not work.

Contributions.

1. **An environment-level null for RL.** Rather than corrupting the model or the labels, we corrupt the environment’s *observation channel*, leaving dynamics and reward intact. The agent trains normally, on a real objective, with nothing to condition on, and this requires no access to task internals (Section 3).
2. **A worked case** of an uninformative explanation passing every conventional check: stable across seeds, separated beyond its own estimation error, and legible (Section 5.1).
3. **Attribution is not identified by behaviour.** An auxiliary training penalty drives the market agent’s dominant feature from 39.9% to 3.2% of attribution mass at no measurable cost in return; the same intervention on a load-bearing feature costs 81% of return. No deception and no attack on the explainer is involved (Section 5.6).
4. **The established null fails in RL, and we show why.** Parameter randomisation never exceeds $z = 2.45$ across sixteen checkpoints, including on a converged CartPole policy scoring 404 of 500, while the environment null reaches $z = 117$. Its width equals the standard deviation of random-initialisation return to three significant figures, so it measures the variance of initialisation (Section 5.5).

5. **Validation with ground truth, synthetic and real.** We construct environments where the correct attribution is known by construction (Sections 4.1 and 5.2), and a positive control on real market episodes where only the objective varies (Section 5.10).
6. **An audit of our own null construction**, which our thesis implies must matter and does (Section 5.11).

What we do not claim. The argument that RL attribution should target outcomes rather than policy outputs is due to Beechey et al. [3], whose SVERL framework unifies behaviour, outcomes and prediction, and whose FastSVERL [2] addresses the scalability, off-policy and temporal problems we only name. Our test adapts the data-randomisation test of Adebayo et al. [1]. Deletion-based fidelity evaluation is established practice in XRL [5, 9]. Our contribution is the RL construction, the demonstration that the incumbent alternative fails, and the empirical consequences.

2 Related work

Shapley values in RL. Beechey et al. [3] give a theoretical account of Shapley-value explanation for RL, showing that earlier applications explained the wrong object and proposing characteristic functions built on agent performance. Their framework identifies three explanatory targets: behaviour, outcomes, and prediction. FastSVERL [2] makes the approach tractable and addresses temporal dependence and off-policy data. We adopt the outcomes framing and implement it independently; we do not extend the theory.

Sanity checks. Adebayo et al. [1] introduced parameter- and data-randomisation tests; Binder et al. [4] examine their limitations. Huber et al. [7] apply the parameter test to perturbation-based saliency for Atari agents and note the data test lacks an RL analogue. We propose one, and show the parameter test is the weaker of the two in this setting.

Manipulating explanations. Slack et al. [12] construct adversarial classifiers that detect the off-manifold points LIME and SHAP query and behave differently on them, yielding innocuous explanations for biased models. That is an attack on the explainer’s sampling. Our steering result differs in kind: we train a normal agent with an invariance penalty, the resulting explanation is *correct* about the resulting model, and the point is that behaviour does not pin the explanation down. It is a non-identifiability result rather than an attack.

Uncertainty in Shapley estimation. Goldwasser and Hooker [6] certify top- k Shapley rankings against

Monte Carlo error. These guarantees concern the estimator. Section 5.4 shows they are orthogonal to whether the explained model learned anything.

Evaluation in XRL. Surveys note that the absence of ground-truth explanations is a central obstacle, and that perturbation-based fidelity is the common recourse [5, 9]. We use deletion curves in that established sense and do not claim them as a contribution.

3 Method

3.1 The statistic

Let $v(S)$ be the expected episode return of an agent observing only the features in $S \subseteq N$, with features outside S replaced at every timestep by draws from a reference distribution. Define the *attribution span*

$$\Delta = v(N) - v(\emptyset). \quad (1)$$

Shapley values [11] and integrated gradients [13] satisfy additivity identities (efficiency and completeness respectively) under which per-feature attributions sum to Δ . The span can therefore be read off either.

It is worth stating precisely what Δ is. It is *not* a Shapley quantity: it is the difference between two masked rollouts, and depends only on the masking scheme. Efficiency tells us the Shapley values happen to sum to it. This has a practical consequence we exploit: Δ costs two evaluations rather than $2^{|N|}$.

3.2 The null


The test compares an observed Δ against the distribution of Δ for agents that had nothing to learn. We construct these by corrupting the observation channel (Figure 1).

Blind-observation null. *Replace every observation with an independent draw from a fixed reference distribution matched to the environment’s observation moments. Leave dynamics and reward untouched. Train normally.*

Such an agent faces the real task and receives real return; it simply cannot condition on anything. This is the RL counterpart of training on permuted labels, and unlike parameter randomisation it yields a *normally trained* policy: the failure mode that occurs in deployment rather than an artificial one.


Matching the reference moments matters. A network fed out-of-range inputs fails for reasons of scale rather than information, which would be a different experiment.

(a) parameter randomisation



weights resampled, so the policy never learned anything.
masking its inputs moves return arbitrarily.

(b) observation corruption (ours)



the policy is trained normally on a real reward.
the only thing withheld is information.

Figure 1: Two ways to build a reference distribution of explanations for an agent that learned nothing. The established check (a) randomises the policy’s parameters, leaving a network that was never trained. Ours (b) leaves the policy and its training untouched and corrupts the environment’s observation channel, so the agent optimises a real reward with nothing to condition on. Only (b) yields a normally trained policy, which is the failure mode that occurs in deployment; Section 5.5 shows the difference is worth 42× in the width of the resulting reference distribution.

3.3 The test

Given Δ_{obs} and null draws $\Delta_1, \dots, \Delta_n$ with mean $\bar{\Delta}$ and standard deviation s , we report

$$p_{\text{rank}} = \frac{\#\{i : |\Delta_i - \bar{\Delta}| \geq |\Delta_{\text{obs}} - \bar{\Delta}|\} + 1}{n + 1}, \quad (2)$$

$$z = \frac{\Delta_{\text{obs}} - \bar{\Delta}}{s}, \quad (3)$$

and require both to reject.

Two statistics are reported because neither suffices. The rank statistic assumes no distributional shape but bottoms out at $1/(n + 1)$: with $n = 8$ it cannot reject at $\alpha = 0.05$ however extreme the observation, and in an early run it reported $p = 0.111$ for a signal lying 10.5 standard deviations outside the null. The normal statistic has resolution but assumes a shape the null need not have. Requiring both keeps the test conservative in the direction that matters.

Degenerate nulls. If every null draw is identical, $s = 0$ and z is undefined. This is not an absence of information but its opposite: any deviation is maximally surprising. We return $z = \pm\infty$ and surface the degeneracy explicitly. It occurs in practice (Sections 5.5 and 5.7) whenever a task lets a blind agent converge on a single behaviour.

4 Experimental setup

4.1 Environments with known answers

Explanation evaluation lacks ground truth [5]. We construct three synthetic corpora in which the correct answer is fixed by construction. Episodes are fixed-length with no early termination, so the horizon is controlled exactly.

- **Planted signal.** One of 18 features is informative

about the binary outcome; the rest are noise or constants.

- **Null.** Identical shapes and frictions, no informative feature. The optimal policy is to abstain.
- **Decoy.** One feature correlates 0.989 with the outcome in the first 60% of episodes and 0.006 in the remainder.

The claim about the planted corpus needs care. The planted feature is the only one informative about the *outcome*. Others can still affect the *return* without predicting anything, because return depends on behaviour: an agent that cannot tell where it is in an episode acts incoherently and pays transaction costs. The test is that the planted feature ranks first with the majority of attribution mass, not that everything else is zero.

4.2 A market with terminal ground truth

Our main empirical setting is a binary contract settling to \$1.00 if BTC is higher after 15 minutes and \$0.00 otherwise, on a regulated prediction market. We use 6,425 settled contracts over 68 days (89,950 transitions, outcome rate 0.4999), with a walk-forward split by time and purge gaps; the test split is evaluated once.

This instrument was chosen for one property: every episode resolves to a known value, so a critic’s $V(s)$ can be checked against realised settlement frequency. Agents are trained with PPO [10]. State is 18 causal features. Actions are target inventory in $\{-1, 0, +1\} \times q_{\text{max}}$. Reward is the change in mark-to-market equity net of costs, which telescopes exactly to net episode P&L.

Transaction costs use the exchange’s published taker schedule, $[0.07C P(1-P)]$, maximised at $P = 0.5$ where this contract trades by construction. With the measured one-cent modal spread this is a 9% round-trip hurdle on a 50-cent notional, and it determines most of what

follows.

4.3 Control tasks

CartPole-v1 (4 features), Acrobot-v1 (6), MountainCar-v0 (2) and Pendulum-v1 (3) provide settings with no relationship to trading and well-understood optimal policies. All four have observation spaces small enough to enumerate every one of the $2^{|N|}$ coalitions, so per-feature Shapley values there are exact. Pendulum’s action space is continuous; we discretise it into nine evenly spaced values so the same categorical PPO trains on every task, at the cost of fine control near the optimum.

4.4 Reproducibility

All results derive from public endpoints with no API keys. Hyperparameters were fixed a priori from synthetic sweeps where the correct answer is known analytically, not tuned on validation performance. The full pipeline is one command; 235 tests cover the environment, the attribution implementations, and six no-lookahead properties verified by mutation testing. Every numerical claim in this paper is asserted against its source artifact by a script that runs as the pipeline’s final stage.

Code and compute. Code, data-ingest scripts and the exact commands that produce every figure and table are at <https://github.com/Aaryansi/nano-11>. All experiments run on CPU: an 8-core Apple M1 with 8 GB of memory, no GPU at any stage. The complete pipeline takes roughly two hours from a cold cache, of which about 27 minutes is rate-limited market ingest, and roughly one hour warm. Policies are two-layer MLPs of 64 hidden units; the 18-feature market agent has 5,636 parameters. The headline experiments are therefore reproducible on a laptop, which is a deliberate property of the setting rather than an accident: the claims concern whether an explanation carries information, and that question does not require scale to pose.

Hyperparameters. PPO with clipped surrogate, GAE, and $\gamma = 1.0$; γ is unity because episodes are finite-horizon and reward telescopes exactly to terminal P&L, so discounting would bias the objective away from the quantity of interest. Market agents train for 100 updates over 5 seeds; explanation-time agents for 60 updates; null draws use $n = 24$ for the main test and 12–16 for the environment and robustness experiments. Full values, including network widths, learning rate, clip range and entropy coefficient, are in the configuration files in the repository rather than restated here.

The exchange serves a moving window, so re-running the ingest returns a different set of contracts. A cold rebuild five days after the original produced 6,425 episodes against 6,428, sharing 5,960 tickers; *overlapping episodes are byte-identical on every field*. Ingest is deterministic per contract and what changes is membership. All con-

Table 1: Test split, 1,284 held-out episodes. p -values are paired bootstrap against always-flat on the same episodes, necessary because per-episode P&L has a standard deviation near 50 against differences near 1.

Policy	Mean P&L	Trades/ep	p vs flat
always-flat	+0.000	0.00	n/a
buy-and-hold	−1.608	1.00	0.23
PPO (5 seeds)	−0.595 ± 0.176	1.37	0.10
logistic (refit)	−5.729	2.52	<0.0001
mean-reversion	−11.615	4.94	<0.0001
random	−18.532	9.33	<0.0001

Table 2: Null-model test on the market. Null from 24 blind-trained agents, span $+8.19 \pm 5.11$.

Case	Span	z	Verdict
planted signal	+70.89	+12.27	informative
null corpus	+11.86	+0.72	not distinguishable
real market	+9.37	+0.23	not distinguishable

clusions reproduced (Section 5.9).

5 Results

5.1 An explanation that passes every check and says nothing

The market is empirically efficient: its implied probability has a weighted mean absolute calibration error of 0.0172 across ten bins, and spot return correlates 0.476 with the outcome but only 0.016 with the residual after the price is accounted for. No policy beats abstention (Table 1); losses track turnover.

PPO’s shortfall is its own. On synthetic data containing no signal by construction the same agent scores -0.47 ± 0.47 ; here it scores -0.595 ± 0.176 . Nothing is left to attribute to the market.

Yet its explanations look entirely reasonable. Across five seeds whose pairwise performance differences are statistically indistinguishable in all ten pairs: behaviour attributions have cross-seed rank correlation 0.849 (min 0.761) with 100% agreement on the top feature and 100% sign agreement; the top-5 ranking is fully certified, all four adjacent gaps exceeding their combined standard errors; and the ranking is semantically legible, time-to-expiry dominating, which reads as an agent that learned when not to trade.

The null test disagrees (Table 2, Figure 2).

The test has power and specificity. The market agent’s explanation sits at the centre of the distribution of explanations of nothing.

We note one consequence for our own reporting. An earlier draft stated that “the whole observation is worth $+5.914$ per episode against being blind” and presented it as a finding. Blind agents produce $+8.19 \pm 5.11$. It was noise.




Figure 2: Attribution spans against the null distribution. The shaded band is the range over 24 agents trained where the observation channel carries no information. The market agent sits at its centre.




Figure 3: A feature that drives the policy but earns nothing, on held-out episodes where it is provably uninformative.

5.2 Validation against ground truth

On the planted corpus, outcome attribution ranks the planted feature first with 55% of total attribution mass. This check is unavailable on real data and is what licenses the rest.

On the decoy corpus, an agent trained where the decoy predicts the outcome (in-sample +47.24, held out -6.83) is explained on held-out episodes where the decoy is provably worthless (Figure 3). Per-decision attribution of $\pi(a | s)$ gives it 6.5% of mass, rank 3 of 18; outcome-based attribution gives it 1.1%, rank 10.

Both statements are true. The policy does key on the decoy, so per-decision attribution is correct about behaviour and misleading about value. An auditor asking what an agent relies on that does not work gets the wrong answer from it.

Deletion curves agree: removing features in the order each ranking considers important, the outcome ranking degrades return faster (AUC -187.8 against -165.3; a random ranking gives +366.6).

5.3 How much edge is required

Sweeping planted-signal strength calibrates the test (Figure 4). Detection begins near +3.45 per episode of *measured* edge. The market agent earns -0.595: an order of magnitude below threshold and of the wrong sign. These edges are in-sample and therefore optimistic; the thresh-




Figure 4: Detection z against the agent's measured edge. The shaded band is where an explanation is not distinguishable from an explanation of nothing.

Table 3: Null construction on four control tasks (CartPole-v1, Acrobot-v1, MountainCar-v0, Pendulum-v1). Entries are mean $\pm$ sd of the span across draws; the sd is what sets the test's power.

Environment	Parameter	Environment	Ratio
CartPole	+55.34 $\pm$ 138.92	-1.35 $\pm$ 3.29	42×
Acrobot	+64.03 $\pm$ 124.94	+0.00 $\pm$ 0.00	degenerate
MountainCar	+0.00 $\pm$ 0.00	+0.00 $\pm$ 0.00	0×
Pendulum	-152.16 $\pm$ 131.42	+11.12 $\pm$ 37.79	3×

old is a lower bound.

5.4 Estimation guarantees do not substitute

Goldwasser and Hooker [6] certify that a top- k ranking is stable under Monte Carlo noise: a guarantee about the estimator. The market agent's top-5, with standard errors near 0.001 on a leading value of 0.26, has all four adjacent pairs separated beyond their combined error and is fully certified by that criterion. The same explanation fails the null test at $z = +0.23$.

A ranking can be certified stable while the explanation it ranks is indistinguishable from an explanation of nothing. Only one of these two properties is commonly reported.

5.5 The established null does not transfer to RL

We construct both nulls for the same task (Table 3, Figure 5). Shapley values here are exact.

A randomly initialised policy behaves arbitrarily, so masking its inputs moves return unpredictably and the reference distribution is very wide. The resulting test has almost no power. Across sixteen checkpoints the parameter null never exceeds $z = 2.45$, including on a converged CartPole policy scoring 404 of 500, while the environment null reaches $z = +117.6$. The two disagree on five of the sixteen. The choice of null is not a detail; it decides the verdict, and the incumbent construction




Figure 5: Spread of each null construction on four control tasks. Width determines power: a very wide reference distribution cannot reject anything. MountainCar is degenerate under both constructions, because no policy in the budget ever reaches the goal.

is the one that fails.

Two environments qualify the claim rather than extend it. MountainCar is degenerate under *both* constructions, at 0.00 ± 0.00 : within our step budget no policy ever reaches the goal, every episode returns -200 , and there is genuinely nothing to explain. The test declines, correctly, and contributes no evidence either way. Pendulum narrows the gap to $3\times$ because its environment null is itself wide (37.79): a blind agent on a dense-reward continuous task still varies in return. The $42\times$ figure is the extreme case, not the typical one.

Acrobot additionally shows the test tracking *whether there is anything to explain* rather than agent quality. Its agent fails completely until it learns (-500 at 10%, 25%, 50% of training), the span is exactly 0.00 there, and the test correctly declines. CartPole behaves differently because its observations matter at every skill level. We do not claim undertrained agents generally produce empty explanations; they do not.

Why the parameter null is wide. The gap has a mechanism, and it is not about explanations at all. Recording each random network’s *unmasked* return alongside its span (Figure 6) shows the parameter null’s width is, to three significant figures, the spread of random-initialisation return: 138.38 against 138.92 on CartPole, 123.04 against 124.94 on Acrobot, 135.49 against 131.42 on Pendulum, and 0.00 against 0.00 on

MountainCar.

The reason is visible in the decomposition $\Delta = v(N) - v(\emptyset)$. For a randomly initialised network the masked term is nearly constant across draws (its standard deviation is 1.8% of the unmasked term’s on CartPole and 5.4% on Acrobot) because a policy that was not using its observations loses almost nothing when they are removed. The span therefore inherits the variance of the unmasked term entire. *Parameter randomisation measures the variance of initialisation, not anything about the explanation being tested*, which is why its reference distribution is wide wherever random networks vary and why a test built on it has no power. An environment null does not have this problem because its agents are trained to convergence and land near a common blind-optimal return.

This supersedes the conjecture we offered in an earlier draft, that the width depended on how often a random network is accidentally competent. That is roughly right but imprecise: what matters is not competence but *variance* in competence, and the relationship is an identity rather than a tendency. Four environments is too few to fit anything, and we report the correspondence rather than a correlation coefficient for that reason.




Figure 6: The parameter null’s width against the spread of random-initialisation return, one point per environment. The points lie on $y = x$: the established null’s reference distribution is the variance of initialisation.

Table 4: Steering one feature’s attribution.

Corpus	Attribution	Return	p
market	39.9% → 3.2%	−3.49 → −1.65	0.66
planted signal	46.5% → 1.5%	+45.04 → +7.92	<0.001

5.6 Attribution is not identified by behaviour

We add an auxiliary penalty during training on the divergence between $\pi(\cdot | s)$ and $\pi(\cdot | s')$, where s' has one feature resampled from the batch marginal: the same interventional perturbation the attribution measures (Table 4, Figure 7).

The attribution is equally suppressible in both cases (92% and 97%). What differs is the price: where the feature is the planted signal, suppressing it destroys 81% of return; on the market it costs nothing measurable. Explanations are steerable at no cost exactly where they are not tracking anything, corroborating Section 5.1 by a route sharing no machinery with the null test. The control matters: had steering succeeded on both corpora, the penalty would merely be defeating the attribution method.

5.7 What the verdict does and does not depend on

Attribution family. Integrated gradients [13], sharing no machinery with Shapley, ranks features at 0.981 correlation with Shapley across five seeds and is comparably seed-stable (0.750 against 0.800). Per-feature attributions are not a Shapley artefact. IG cannot test the outcome-level claim: episode return is not differentiable through the environment, so IG has no outcome analogue. Its behaviour-level span reports the market agent as informative ($z = +4.58$), which is consistent rather than contradictory and is precisely the behaviour-

versus-outcome distinction of Beechey et al. [3]. The agent’s behaviour does respond to its observations; that responsiveness earns nothing.

Credit assignment. Leave-one-out ($v(N) - v(N \setminus i)$) and only-one-in ($v(\{i\}) - v(\emptyset)$) are the extremes of the average Shapley takes, and neither satisfies efficiency, so each carries its own total. On the planted signal all three fire ($z = +10.66, +10.96, +2.29$); on the market none does ($z = -0.42, -1.57, -0.72$). The verdict is scheme-independent. *The per-feature ranking is not:* leave-one-out and only-one-in rank features at only 0.309 correlation with each other on the same agent. Whether there is anything to explain is robust; which features matter is not, and a ranking from a single scheme should not be reported as the answer.

Horizon. We hypothesised that the parameter null’s excess variance arises because randomising weights perturbs the measurement’s *domain*, return being a functional of the policy-induced state distribution, and that this compounds along the trajectory. The prediction is that its variance grows with horizon faster than the environment null’s. It does not: over horizons 4 to 28 both grow near-linearly at indistinguishable rates ($\alpha = 0.94$, $r^2 = 0.995$ against $\alpha = 0.83$, $r^2 = 0.998$), the ratio flat at 1.7–1.8 $\times$. The hypothesis is rejected. The gap is instead accounted for by the initialisation-variance identity of Section 5.5, which involves no compounding along the trajectory and is consistent with the flat ratio found here.

The experiment did establish something else. At horizon 56 the environment null collapses to a standard deviation of 0.04 with two distinct values across sixteen agents: every blind agent converged on the same abstaining policy. This is the Acrobot degeneracy again, and it now has a mechanism: the environment null degenerates when the task affords enough time to learn that inaction is optimal.

5.8 Off-manifold masking

The standard objection to perturbation attribution applies here and deserves a direct answer. We define $v(S)$ by replacing the features outside S with draws that ignore the features inside S . Our 18 market features are correlated, so those draws break the joint distribution and the policy is scored on states that never occur. Slack et al. [12] build their attack on precisely that gap. If the span reflected distribution shift rather than information, the headline result would be an artefact of the masking.

The span cannot be an off-manifold artefact. This is structural rather than empirical, and we state it as such. The statistic is $\Delta = v(N) - v(\emptyset)$, and neither coalition evaluates an off-manifold state. $v(N)$ masks nothing. At $v(\emptyset)$ our implementation replaces the entire observation from a *single* reference row, so the syn-





Figure 7: Attribution is equally suppressible on both corpora (left). Only where the feature is load-bearing does suppressing it cost return (right).




Figure 8: Distance from a masked state to the nearest real state, in per-feature standard deviations, against the number of features replaced. The dashed line is where real held-out states sit relative to the reference sample. Marginal masking moves *above* that floor as more features are replaced; neighbour-conditioned masking stays at or below it. The fully-masked point is omitted: those states are drawn from the reference sample itself, so their distance to it is zero by construction rather than by evidence.

thetic state is itself a real observation, merely a different one. Conditioning is vacuous on an empty kept set, since $p(x \mid \emptyset) = p(x)$. A conditional masker must therefore agree with the marginal one on Δ exactly, and ours does: +8.838 under both. That equality confirms the implementation respects the argument; it is not independent evidence for it.

The concern is real in the interior. It would be convenient but false to stop there, because the argument above says nothing about intermediate coalitions, which is where per-feature attributions live. We implemented a conditional masker that draws replacements from the $k = 16$ reference rows nearest the current state in the kept features only, approximating $p(x_{\bar{S}} \mid x_S)$. Measur-

ing the distance from masked states to the nearest real state (Figure 8) shows the two modes are not equivalent: with 14 of 18 features replaced, marginal masking sits at 0.486 against a real-state floor of 0.381, while conditional masking sits at 0.186. Marginal masking does leave the manifold, by a measurable amount, in the interior.

Consequences. Re-deriving the verdict entirely under conditional masking, against a null of 12 conditionally-masked blind agents, reproduces it: the planted signal fires at $z = +9.61$ and the real market does not at $z = -0.55$. The per-feature ranking is another matter. Leave-one-out values under the two modes correlate at only 0.767 by rank, and *they disagree on which feature is most important*. This is the same split Section 5.7 found for credit assignment: whether there is anything to explain is robust, which feature to name is not. A practitioner may report the verdict; the ranking should carry the masking mode with it.

5.9 Reproduction from scratch

All figures above come from a corpus rebuilt after deleting every cached file. An earlier corpus, collected five days before and sharing 93% of its contracts, gave PPO -0.418 ± 0.260 ($p = 0.129$) against -0.595 ± 0.176 ($p = 0.100$), and null-test z of $+12.27 / +0.72 / -0.15$ against $+12.27 / +0.72 / +0.23$. Every conclusion is identical on both. Buy-and-hold moved from $+0.068$ to -1.608 , confirming its earlier positive figure reflected that period’s 52.3% outcome rate rather than an edge.

5.10 A positive control on real data

Every case where the test fires elsewhere in this paper is synthetic, which invites the obvious objection: perhaps it detects planted signal rather than information, and would decline on any real corpus. The control holds the corpus, the 18 features, the normalizer, the walk-forward split and the null construction fixed, and varies only the

objective.

The market’s implied probability is well calibrated, with a weighted mean absolute error of 0.0172. A calibrated price is highly informative *about the outcome* while offering nothing to trade against, because the fee schedule takes more than the edge. So we score the same agent architecture on the same real episodes for *calling* the settlement rather than trading it: +1 for each step the call agrees with the eventual outcome, −1 when it disagrees, 0 for abstaining. Reward may reference the settlement; an observation may not, and none does. The environment subclasses the trading one, so the observation construction and normalizer handling are literally the same code and a difference in results cannot be a difference in the environment.

Against a null of 24 agents trained on these same real episodes with the observation channel blinded, the prediction agent’s span is +7.46 against a null of -0.48 ± 2.34 , giving $z = +3.40$: informative, and unchanged in every one of 4,000 resamples. The trading agent on the identical episodes, against the identical null construction, gives +7.41 against $+2.47 \pm 2.29$ and $z = +2.16$, which the two-part rule declines. Note that the two spans are almost the same size; what separates them is the null. A blind agent cannot call the outcome at all, so the prediction null sits at zero, while a blind agent can still avoid trading and so retains most of the trading agent’s advantage.

The same feature stream therefore carries usable information for one objective and not for the other, and the test says so. This also sharpens what the negative result means. It is not that the market data is uninformative; a calibrated price is very informative. It is that the information is not *monetisable* through a fee schedule that charges 9% round trip, and an attribution of the trading agent cannot tell those two situations apart.

5.11 Our own null construction, checked

Section 5.5 argues that the choice of null decides the verdict and that practitioners do not check it. That argument applies to this paper, so we checked.

Section 3 defines the null as corrupting the observation channel while leaving dynamics and reward intact. The null used for the market result does something slightly different: its agents are trained *and* attributed on synthetic signal-free corpora, so the corpus changes as well as the information, and their spans are measured on synthetic episodes while the observed span is measured on real ones. The two quantities being compared are not measured on the same thing.

We therefore rebuilt the null the way the method section defines it: agents trained on the *real* training episodes with observations replaced by moment-matched noise, and attributed on the same real test episodes as the observation. Holding the agent fixed at the check-point the paper explains and the measurement corpus fixed at the real test split, the null construction is then




Figure 9: One agent, one observed span, two null constructions. The vertical line does not move; only the reference distribution does. The synthetic-corpus null is nearly twice as wide, because it absorbs corpus-to-corpus variation the observation does not have.

the only thing that varies (Figure 9, $n = 32$ each):

Null	Mean $\pm$ sd	z	Verdict
synthetic corpora	$+9.07 \pm 4.91$	-0.34	not distinguishable
blinded real	$+2.26 \pm 2.21$	$+2.33$	not distinguishable

The verdict survives, but not comfortably, and the reason it survives is worth stating. Under the blinded-real null the normal statistic *does* reject at $z = +2.33$; the rank statistic does not, and Section 3 requires both. The two-part rule is doing real work here rather than being a formality.

Resampling each null (Section 5.12) shows how much weight the two comparisons carry. The synthetic-corpus verdict is unchanged in 100% of replicates; the blinded-real verdict in 65%. We report the market result as undecided under the stricter construction rather than as confirmed, and we recommend the blinded-real form for future use: it holds the corpus fixed, which is what isolates information as the variable.

An earlier run of this comparison at $n = 12$ put the two constructions on opposite sides of the threshold. That disagreement was a small-sample artefact and disappeared at $n = 32$. We mention it because it is the failure mode this paper is about, encountered in our own results, and because $n = 12$ null agents is a common budget.

5.12 Confidence intervals on the verdicts

Every z in this paper divides by a standard deviation estimated from 12 to 32 null agents. Reporting it to two decimals hides that. We resample the null draws with replacement and re-apply the full two-part decision rule to each replicate, which gives a percentile interval for z and, more usefully, the fraction of replicates in which the verdict is unchanged.

Every verdict in the paper is stable in 100% of replicates except the two blinded-real market comparisons,

at 64% and 65%. The planted signal’s interval is $[-10.3, +16.9]$ and the market’s, under the null used in Section 5.1, is $[-0.16, +0.73]$: nowhere near the threshold in either direction. This bootstraps the null only. Uncertainty in the observed statistic, a mean over held-out episodes, is not included, so these intervals are optimistic in a way we would rather state than hide.

5.13 Implications for oversight

The case for interpretability as an oversight mechanism is that seeing which inputs drive a policy tells us something about whether to trust it. Three of our results bear on that argument, and none of them requires an adversary.

Passing every check is not evidence. On a corpus where the correct answer is known to be “nothing”, an agent’s attribution is stable across seeds, separated beyond its own estimation error, and semantically legible. Every property a practitioner would check is present and the explanation is empty. Stability in particular is worth naming: it is the most commonly reported evidence that an explanation is trustworthy, and Section 5.14 records that we expected instability to be the tell and were wrong. Emptiness and instability are unrelated.

Explanations are cheap to forge, without deception. Section 5.6 drives a feature’s attribution from 39.9% to 3.2% of total mass at no measurable cost in return. This needs no deceptive policy, no attack on the explainer, and no knowledge of who will audit the model: an auxiliary training term is enough, and the resulting explanation is *correct* about the resulting agent. The cost of forgery is exactly the importance of the feature, which is zero for features that do not matter. An oversight regime that reads attributions as evidence of what a model relies on is therefore reading a quantity the developer can set, in the region where it would most want a guarantee.

The verdict depends on a choice nobody reports. Section 5.5 shows the established sanity check has almost no power in RL, and Section 5.11 shows that our own headline verdict moves when the null construction is corrected. A reported attribution is not interpretable without the reference distribution it was measured against, and that reference is currently neither standardised nor, in most papers, stated.

What we would ask for. Report the null construction alongside any attribution offered as evidence. Treat a per-feature ranking as a joint property of the attribution scheme and the masking distribution, since Sections 5.7 and 5.8 show it moves under both, to the point of disagreeing on the most important feature. The aggregate verdict, whether there is anything to explain at

all, is the part that survives; it is also the weaker claim, and it is the one we would trust.

5.14 Hypotheses we rejected

Four predictions in this project turned out to be wrong. They are collected here rather than left implicit in the sections that tested them, because a reader assessing whether the surviving claims were stress-tested should be able to see what did not survive.

One further claim is stated in the paper as structural rather than empirical. The span’s invariance to the masking mode (Section 5.8) follows from the construction, since $p(x \mid \emptyset) = p(x)$; the measured equality confirms the implementation respects that argument and is not independent evidence for it. We separate the two because a numerical agreement that must hold is easily mistaken for one that might not have.

6 Limitations

The construction is adapted. The test is the data-randomisation test of Adebayo et al. [1] with an RL-appropriate null; the outcomes framing is Beechey et al. [3]’s; deletion-based fidelity is standard.

The span measures information in aggregate. An explanation could clear the test and still misattribute among features. Section 5.7 shows the ranking is scheme-dependent even when the verdict is not, and Section 5.8 shows it is masking-dependent too, to the point of disagreeing on the most important feature. Only the verdict survives both.

Conditional masking is approximated, not solved. Our conditional masker is k -nearest-neighbour in the kept features, which is a practical proxy for $p(x_{\bar{S}} \mid x_S)$ and inherits the usual difficulties in high dimension. It is adequate to show the two modes differ; it is not a claim to have sampled the true conditional. Neither mode fixes the deeper issue that replacements are redrawn each timestep, so masked *states* are plausible while masked *trajectories* are not.

Temporal credit assignment is named, not solved. Masking a feature for a whole episode measures whether it mattered, not when. This is FastSVERL’s [2] territory.

Off-policy attribution is not handled. Masked coalitions induce a different policy and therefore a different state distribution; estimates are on-policy for the masked policy, not counterfactuals about the unmasked one.

The width identity is measured on four environments. The correspondence between the parameter null’s width and initialisation-return variance is close to exact on each task we ran, and it has a mechanism, but four points cannot establish a law. We report the correspondence, not a fitted relationship.

Scope. One real domain and four classic-control tasks; one attribution family for the outcome-level tests.

Table 5: Predictions this project made and then rejected on its own evidence. None of the surviving results depends on any of them.

Hypothesis	Prediction	Outcome
An agent with no edge produces unstable explanations, and instability is the tell.	Low cross-seed rank correlation.	Rejected. 0.850 mean rank correlation, unanimous on the top feature. Emptiness and instability are unrelated, which is what makes the empty case hard to catch.
Outcome-level attribution, being tied to return, is more stable than behaviour-level attribution.	Higher cross-seed rank correlation for outcomes.	Rejected, and reversed. Outcomes reach 0.347 against behaviour’s 0.850. Attributing to return is noisier, not steadier.
The parameter null is wide because randomising weights perturbs the measurement’s domain and the error compounds along the trajectory.	Its variance grows with horizon faster than the environment null’s.	Rejected. Over horizons 4 to 28 both grow at indistinguishable rates ($\alpha = 0.94$ against 0.83), ratio flat at $1.7\text{--}1.8\times$.
The parameter null is wide because a randomly initialised network is sometimes an accidentally competent policy.	Width tracks how often random policies do well.	Superseded. Roughly right but imprecise: the width is not related to competence but <i>equal</i> to the standard deviation of initialisation return (Section 5.5). A tendency was replaced by an identity.

All four control tasks have small observation spaces, so nothing here speaks to pixel-based or high-dimensional agents, where the reference distribution is itself the hard problem. Box2D environments were deliberately excluded: they require a system-level build dependency, which would trade the one-command reproduction for one more observation dimension. Corpus membership is not reproducible across time, though content is.

7 Conclusion

An attribution can be stable, precisely estimated, plausible, and empty. The checks practitioners apply (consistency across runs, separation beyond estimation error, semantic legibility) do not detect this, because none of them ask whether the explained agent learned anything. A null-model test does, using a statistic the attribution framework already provides. The established test for this purpose does not survive the move to reinforcement learning, and we recommend against its use there without the measurement in Section 5.5.

References

- [1] Julius Adebayo, Justin Gilmer, Michael Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. Sanity checks for saliency maps. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018. arXiv:1810.03292.
- [2] Daniel Beechey and Özgür Şimşek. Approximating Shapley explanations in reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2511.06094.
- [3] Daniel Beechey, Thomas M. S. Smith, and Özgür Şimşek. Explaining reinforcement learning with Shapley values. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. arXiv:2306.05810.
- [4] Alexander Binder, Leander Weber, Sebastian Lapuschkin, Grégoire Montavon, Klaus-Robert Müller, and Wojciech Samek. Shortcomings of top-down randomization-based sanity checks for evaluations of deep neural network explanations. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [5] Zelei Cheng, Jiahao Yu, and Xinyu Xing. A survey on explainable deep reinforcement learning. *arXiv preprint arXiv:2502.06869*, 2025.
- [6] Jeremy Goldwasser and Giles Hooker. Statistical significance of feature importance rankings. *arXiv preprint arXiv:2401.15800*, 2024.
- [7] Tobias Huber, Benedikt Limmer, and Elisabeth André. Benchmarking perturbation-based saliency maps for explaining Atari agents. *Frontiers in Artificial Intelligence*, 5:903875, 2022. doi: 10.3389/frai.2022.903875.
- [8] Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [9] Stephanie Milani, Nicholay Topin, Manuela Veloso, and Fei Fang. Explainable reinforcement learning: A survey and comparative review. *ACM Computing Surveys*, 56(7):1–36, 2024. doi: 10.1145/3616864.
- [10] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal

policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

- [11] Lloyd S. Shapley. A value for n-person games. *Contributions to the Theory of Games*, 2(28):307–317, 1953.
- [12] Dylan Slack, Sophie Hilgard, Emily Jia, Sameer Singh, and Himabindu Lakkaraju. Fooling LIME and SHAP: Adversarial attacks on post hoc explanation methods. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, 2020. arXiv:1911.02508.
- [13] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. Axiomatic attribution for deep networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017.